

Finding Resource Manipulation Bugs with Monitor Automata on the Example of the Linux Kernel

Anders Fischer-Nielsen
afn@itu.dk

2020

Contents

Contents	2
1 Introduction	3
2 Building Upon the EBA Framework	3
2.1 Regions	4
2.2 Effects	4
2.3 Shapes	5
2.4 Shape-and-effect Inference	6
2.5 Effect-CFG Abstraction	7
3 An Algorithm for Finding Bugs	7
3.1 Monitor Templates	7
3.2 Control Flow	11
4 Finding Double-Unlock Bugs in Practice	15
4.1 Integration into EBA	15
4.2 Monitor Signatures	21
5 Evaluation	24
6 Future Work	24
7 Conclusion	24
References	25
8 Appendix	26

1 Introduction

The Linux kernel supports a vast array of computer architectures and runs on a multitude of devices from embedded devices, through personal computers to large servers; on wireless access points, smart TVs, smartphones, refrigerators. Errors in the Linux kernel therefore affect a multitude of devices and can therefore have a potential significant negative impact.

An important aspect of kernel programming is management and manipulation of resources, be it devices, file handles, memory blocks, and locks. Shared-memory concurrency and locks are used extensively in the C source code of the Linux kernel in order to allow parallelization of subsystems within the kernel while at the same time avoiding race conditions. Static analysers allow detection of errors in the C source code of the Linux kernel by reasoning about this resource manipulation. A control flow graph can be found for the components of the kernel, which can then in turn be statically analysed to detect possible resource manipulation errors.

[3]

2 Building Upon the EBA Framework

The EBA framework provides *shape-and-effect inference* and allows extracting an *effect-annotated control flow graph* (effect-CFG) from a Linux kernel program. The annotated control flow graph can be statically analyzed in order to find possible bugs in the given program. Figure 1 shows an example of a bug found in the Linux kernel¹ and the data types provided by the framework. These data types have been used in this thesis to develop new resource manipulation bug checkers. This section will explain the definitions by Abal et. al. [1] which are built upon in this thesis in order to develop new bug checkers.

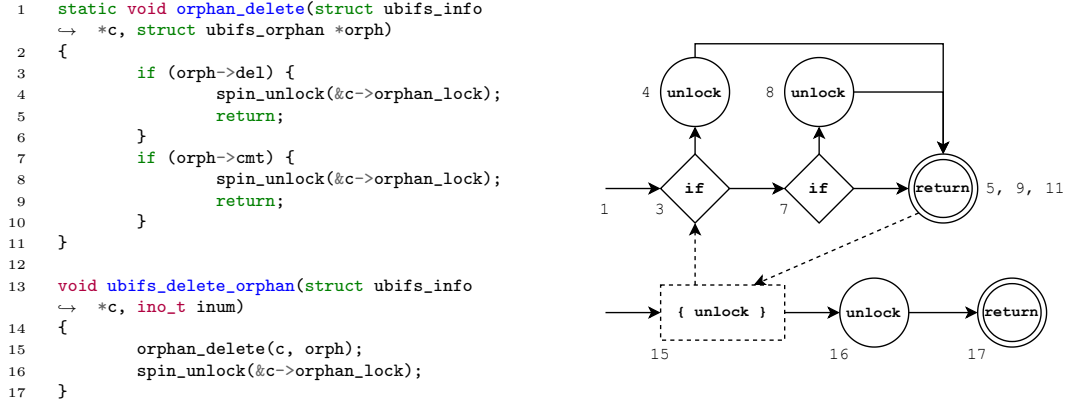


Figure 1: An illustration of a double-unlock bug (4dd75b33) found in the Linux kernel and the analysis data types provided by EBA. The relevant code is shown on the left-hand side and on the right-hand side is the *effect-CFG* with *lock* and *unlock effects* provided by EBA. The numbers next to CFG nodes show the corresponding line numbers.

The EBA framework is built upon in this thesis, though it is generally seen as a black box. An understanding of certain definitions described by Abal et. al. [1] is required to see how we build

¹This bug was patched in commit 4dd75b33.

upon the existing work. Three main definitions are of note here, namely *shapes*, *regions* and *effects*. EBA employs CIL [4] as a lightweight intermediate representation lacking implicit type conversions of the C code of the Linux kernel. Abal et. al define a base type system for a smaller language than CIL. This chapter will provide a summary of the definitions we also use in this thesis to define our work.

Abal et. al. define the base type system of the *Shape-and-Effect System* as

$$\begin{array}{ll} \text{l-value types } T^L & : \quad \text{ref } T^R \quad | \quad \text{ref } (T_1^R \times \dots \times T_n^R \rightarrow T_0^R) \\ \text{r-value types } T^R & : \quad \text{int} \quad | \quad \text{ptr } T^L \end{array}$$

l-value (T^L) types and r-value (T^R) types correspond to the left and right side of assignments in C. A reference type, **ref** T represents a memory cell, holding objects of the type T . For example, **ptr** **ref** T is the current address of the reference for the objects T in memory.

$$\begin{array}{ll} \text{l-value expressions } L & : \quad x \quad | \quad f \quad | \quad *E \\ \text{r-value expressions } E & : \quad n \quad | \quad E_1 + E_2 \quad | \quad \text{if } (E_0) E_1 \text{ else } E_2 \quad | \quad (T)E \\ & | \quad \text{new } x : T = E_1; E_2 \quad | \quad !L \quad | \quad \&L \quad | \quad L_1 := E_2; E_3 \\ & | \quad \text{fun } T f (T_1 x_1, \dots, T_n x_n) = E_1; E_2 \quad | \quad L_0(E_1, \dots, E_n) \end{array}$$

L-value expressions (L) represent memory locations and will always be assigned reference types (T^L). Function values are immutable, while other variables (x) are not. $*E$ represents the dereferencing of a pointer, which is looking up the reference cell in memory, as seen in C. R-value expressions are *values*, such as integers (n) and pointers. $(T)E$ is a cast, as found in C, and will convert the value E to the type T . $\text{new } x : T = E_1; E_2$ represents the introduction of a new variable, x , which is initialized in E_1 and available in E_2 . x is the name of the memory cell where E_1 is stored, and has the type **ref** T . $!L$ will read an l-value, and pointer values pointers can be obtained with $\&L$. $L_1 := E_2; E_3$ allows assigning a different value (E_2) before evaluating E_3 . $\text{fun } T f (T_1 x_1, \dots, T_n x_n) = E_1; E_2$ represents the declaration of a function, f , which will be visible in E_2 , similar to **new**. The function f will bind the parameters x_1, \dots, x_n and evaluate E_1 . Loops and **gotos** are not modelled in this system.

2.1 Regions

Regions are an abstract representation of memory, denoted as ρ . Variables are names for memory cells, and aliasing — when multiple variable names actually point to the same memory — can therefore happen. These possibly aliased memory cells are tracked by the shape-and-effect system as regions. The system will assign a region, ρ , to each reference type in the source code, and attempt to detect aliased variables, by unioning these regions when it can no longer distinguish the regions.

2.2 Effects

Effects represent how expressions affect regions. For example, an expression which reads a memory location will have the effect of reading that region. Likewise, expressions writing to memory locations will have the effect of writing to that region. An example of a set of effects is $\varphi = \{\text{read}_\rho, \text{read}_{\rho'}, \text{write}_{\rho'}\}$, where the region ρ is being read and the region ρ' is being both read and written.

2.3 Shapes

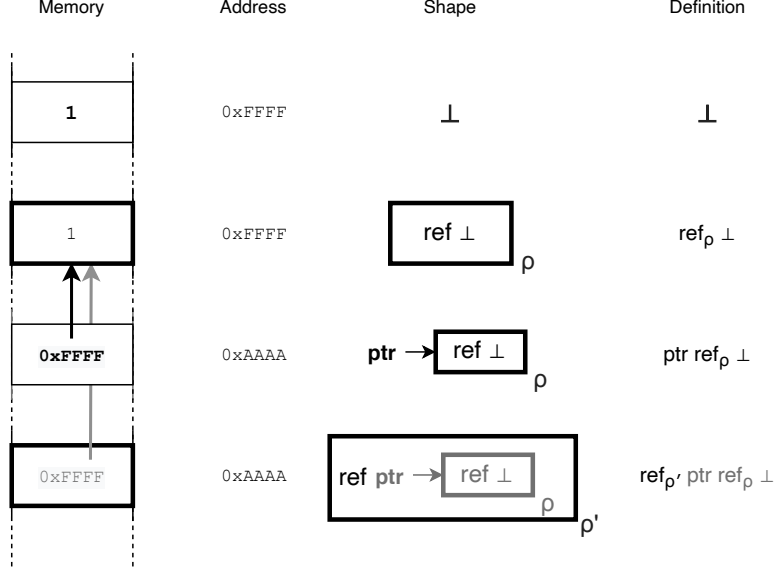


Figure 2: An illustration of *shapes* and how they represent values in memory. Arrows represent pointers, and bold cells refer to the cell itself, not its contents.

A shape approximates the memory representation of an object, and this shape is fixed and kept across type casts. Shapes are annotated with regions showing a "points-to" relationship between references. Abal et. al. defines shapes in the following terms.

$$\begin{aligned}
 \text{l-value shapes } Z^L & : \text{ref}_\rho Z^R \mid \text{ref}_\rho (Z_1^L \times \dots \times Z_n^L \xrightarrow{\varphi} Z_0^R) \\
 \text{r-value shapes } Z^R & : \perp \mid \text{ptr } Z^L \mid \zeta
 \end{aligned}$$

Shapes are also divided into l-value and r-value shapes in a similar fashion as the aforementioned type system, albeit without an integer type and with shape variables, ζ . An r-value is the shape objects where the *atomic* shape, \perp , indicates that an object has no relevant structure — e.g. an integer. A pointer expression has the pointer shape, $\text{ptr } Z^L$, where Z^L is the shape of the target reference cell of the pointer. This means that a pointer represents the address of a reference cell, and therefore a pointer shape necessarily encloses a reference shape. If a pointer is being cast, its integer value will then have a pointer shape, $\text{ptr } Z^L$. ζ represents arbitrary r-value shapes. These definitions and how they represent values in memory are illustrated in Figure 2.

For functions, these function shapes are mapped to *function shapes*,

$$\text{ref}_{\rho_0} \left(Z_1^L \times \dots \times Z_n^L \xrightarrow{\varphi} Z_0^R \right),$$

where the memory region ρ_0 identifies the function and is used by EBA to keep track of calls to the function.

Function shapes represent an abstraction of the shapes given to a function as well as the shape of the result of the function. The parameters $Z_1^L \times \dots \times Z_n^L$ correspond to the shapes of parameters

the function is given and Z_0^R is the shape of the result of the function. Since function parameters are stored in stack variables, these parameters are l-value shapes. The *latent effect*, represented as $\xrightarrow{\varphi}$ above represents the effects that may happen when executing the function. These effects are described in the following section. Function shape schemes along with the correlation between types and shapes are described in more detail by Abal et. al. [1].

An environment Γ maps variables v to their corresponding reference shapes: $\Gamma(v) = \text{ref}_\rho Z$. v is effect variables, described in the following. As already mentioned, function shapes are represent effects that may happen when executing a function, but this *may* can be made more concrete using function inlining, in which case the actual effects of calling a function within another function can be inferred concretely. The use of inlining is detailed in Section 4.

2.4 Shape-and-effect Inference

Abal et. al. present inference rules $\vdash_L \subseteq \text{ENV} \times \text{L-VALUE} \times \text{SHAPE} \times \text{EFFECT}$ for *shape-and-effect inference*, which allows determining what shape a given expressions has as well as determining what the effects of evaluation the expression are. This is expressed as the judgments:

- $\Gamma \vdash_L L : \text{ref}_\rho Z \& \varphi$, specifying that under Γ , the l-value expression L has the shape $\text{ref}_\rho Z$, resulting in the effects φ
- $\Gamma \vdash_E E : Z \& \varphi$, specifying that under Γ , the r-value expression E has the shape Z resulting in the effects φ

Following the previous examples of read and write effects, two examples of inference rules defined by Abal et. al. are shown in the following corresponding to these effects. The FETCH and ASSIGN inference rules allow determining whether an expression will have the effect of reading or writing.

$$\begin{aligned} \text{[FETCH]} \quad & \frac{\Gamma \vdash_L L : \text{ref}_\rho Z \& \varphi}{\Gamma \vdash_E !L : Z \& \varphi \cup \{\text{read}_\rho\}} \\ \text{[ASSIGN]} \quad & \frac{\Gamma \vdash_L L : \text{ref}_\rho Z \& \varphi_1 \quad \Gamma \vdash_E E_1 : Z \& \varphi_2 \quad \Gamma \vdash_E E_2 : Z' \& \varphi_3}{\Gamma \vdash_E L := E_1; E_2 : Z' \& \varphi_1 \cup \varphi_2 \cup \{\text{write}_\rho\} \cup \varphi_3} \end{aligned}$$

[FETCH] will fetch a value in a ref_ρ by the use of the bang operator, $!L$, and give an object of a shape, Z , leading to the effect read_ρ being added to the set of effects, φ . Applying [FETCH] will result in an *l-value*, which under the environment has the shape $\text{ref}_\rho Z$ with the effects φ . [ASSIGN] will evaluate both expressions E_1 and E_2 , resulting in effects for both E_1 , E_2 , along with the write_ρ effect, as well as existing effects being combined into the effects set. The expression E_1 will be written to the location L . The left and right-hand sides of the expression, L and E_1 must be of the same shape, Z , while the result of evaluating E_2 can be of a different shape, Z' .

The full inference system is presented in the work by Abal et. al [1], FETCH and WRITE are merely presented here as examples of effect inference.

Abal et. al. axiomize the behaviour of certain operations, f , with a signature, $Z_i^L \xrightarrow{\varphi} Z_0$. The axiom specifies shapes of the input arguments expected by the function, Z_i^L , the shape of the

output, Z_0 , and the resulting effects, φ . The authors present an example of two axioms for locking and unlocking functions found within the Linux kernel.

$$\begin{aligned}\text{spin_lock} : \quad & \text{ref}_{\rho_1} \text{ptr ref}_{\rho_2} \zeta \xrightarrow{\text{lock}_{\rho_2}} \perp \\ \text{spin_unlock} : \quad & \text{ref}_{\rho_1} \text{ptr ref}_{\rho_2} \zeta \xrightarrow{\text{unlock}_{\rho_2}} \perp\end{aligned}$$

Axioms have been defined by Abal et. al. for multiple functions in the kernel, allowing inferring what the effects are of using these built-ins. These specify that **spin_lock** and **spin_unlock** receive pointers as arguments, ρ_1 , pointing to an object, ρ_2 . The effects above the arrows indicate that the functions **lock** and **unlock** the object in ρ_2 , respectively.

2.5 Effect-CFG Abstraction

Abal et. al present the *Effect-based Control-Flow Graph* (φ -CFG). This is described as a CFG where nodes represent program locations and edges specify the control flow, annotated with variables with their memory shapes, and nodes with the effects inferred for their corresponding locations.

3 An Algorithm for Finding Bugs

In order to detect resource manipulation bugs, I will present an algorithm for detecting these bugs utilizing the definitions presented by Abal et. al [2]. This section will present my definitions of *monitor templates* as finite state machines operating on *effects* on *regions*, allowing the definition of bug behaviour and the detection of such behaviour.

A finite state machine is a tuple $(\Sigma, S, s_0, \delta, F)$, where Σ is an alphabet, S is a finite non-empty set of states, s_0 is an element of S and initial state, δ is the state-transition function $\delta : S \times \Sigma \rightarrow S$ and F is the possibly empty set of final states and a subset of S . I will use such state machines to represent the code under analysis and the properties I wish to detect in input source files.

3.1 Monitor Templates

I use monitor automata to analyze the control flow of a given input source file and detect whether possible bug are present in. A monitor automaton changes state based on what is happening in the control flow of the program. When the automaton reaches an error state, then a possible bug has been discovered. The effect analysis provided by EBA allows monitoring which effects program points have, and monitor automata can then monitor these effects in order to determine whether possible bugs are present.

EBA infers what effects happen on a memory region, which is an abstract variable or value in the heap. Monitor automata track effects happening on a given region in order to determine whether they could be manifesting buggy behaviour. Both effects and regions must be tracked. For example, it is common to have multiple locks on different regions within programs, but a lock on one region followed by a lock on a different region does not necessarily mean that a locking bug is present unless these locks happen consequently on the same region. Monitor templates are defined formally as follows.

Given a region variable ρ , a monitor template is defined as the tuple $X_\rho(\Sigma, S, s_0, \delta, F, E)$ where Σ is the alphabet of effects registered on memory objects represented by a region ρ , S is a finite non-empty set of states, s_0 an initial state and an element of S , δ is the state-transition function $\delta : S \times \Sigma \rightarrow S$, F is the non-empty set of final states and a subset of S , E is the non-empty set of error states, and a subset of F . In other words, $X_\rho(\Sigma, S, s_0, \delta, F)$ is a finite state machine over regions and effects, augmented with the set of error states E . An illustration of such a monitor template can be seen in Figure 3.

From this definition, I distinguish two kinds of monitor templates: *long-term* and *short-term*. *Short-term* templates will monitor effects happening on a region only until a final state or error state is reached. *Long-term* templates operate as *short-term* templates, though they only include a single final state, which is the error state. They will therefore monitor indefinitely until only an error-state is found; no early termination is possible. This distinction is made since early termination of monitors might result in performance improvements in the implementation of these monitors, since it may especially reduce the number of monitors active in the analyzers' memory.




Figure 3: An illustration of a monitor template.

I will present several examples of these monitors in the following sections. Though these are examples, they are important since the implementation of the monitors follows these examples.

Our monitors operate on the set of possible effects of a statement in the Control-flow Graph. EBA allows defining new effects tracking a host of different operation. In this thesis I will only be using the following effects, defined as $E = \{\text{alloc}, \text{free}, \text{read}, \text{write}, \text{uninit}, \text{call}, \text{lock}, \text{unlock}\}$ by Abal [2]. The effects and what they represent can be seen in Table 1.

alloc	The allocation of a memory location
free	The freeing of a memory location
read	The reading of a memory location
write	The writing to a memory location
uninit	Indicating that a memory location is uninitialized
call	The call of a function
lock	The locking of a memory location
unlock	The unlocking of a memory location

Table 1: Effects and what they represent.

Short-term Double-lock Monitor Template




Figure 4: An illustration of a short-term double-lock monitor template.

A double-lock monitor detects two consecutive locks on a memory location with no unlock in between them leading to an infinite spinlock. Given a region ρ , a double-lock monitor template is defined as the tuple $(\Sigma, S, s_0, \delta, E, F)$ where:

- $\Sigma = \{lock_\rho, unlock_\rho\}$, a subset of E
- $S = \{locked_\rho, unlocked_\rho, error_\rho\}$
- $s_0 = unlocked_\rho$
- $\delta =$ the relation $\{(unlocked_\rho, lock_\rho, locked_\rho), (locked_\rho, unlock_\rho, unlocked_\rho), (locked_\rho, lock_\rho, error_\rho), (unlocked_\rho, unlock_\rho, unlocked_\rho)\}$
- $E = \{error_\rho\}$
- $F = \{unlocked_\rho, error_\rho\}$

It is worth noting that this is a *short-term* monitor, as $F \neq E$. This monitor will therefore terminate when encountering a legal use of locks followed by an unlock. An illustration of this monitor template can be seen in Figure 4.

Long-term Double-unlock Monitor Template




Figure 5: An illustration of a long-term double-unlock monitor template.

A double-unlock monitor detects two consecutive unlocks on a memory location with no lock in between them, leading to undefined behaviour. Given a region ρ , a double-unlock monitor template is defined as the tuple $(\Sigma, S, s_0, \delta, E, F)$ where:

- $\Sigma = \{\text{unlock}_\rho, \text{lock}_\rho\}$, a subset of E
- $S = \{\text{locked}_\rho, \text{unlocked}_\rho, \text{error}_\rho\}$
- $s_0 = \text{unlocked}_\rho$
- $\delta = \text{the relation } \{(\text{locked}_\rho, \text{unlock}_\rho, \text{unlocked}_\rho), (\text{locked}_\rho, \text{lock}_\rho, \text{locked}_\rho), (\text{unlocked}_\rho, \text{lock}_\rho, \text{locked}_\rho), (\text{unlocked}_\rho, \text{unlock}_\rho, \text{error}_\rho)\}$
- $E = \{\text{error}_\rho\}$
- $F = E$

It is worth noting that this is a *long-term* monitor, as $F = E$. An illustration of this monitor template can be seen in Figure 5.

Long-term Double-free Monitor Template




Figure 6: An illustration of a long-term double-free monitor template.

A double-free monitor detects two consecutive frees on a memory location with no allocation in between them, potentially leading to modification of unexpected memory locations. Given a region ρ , a double-free monitor template is defined as the tuple $(\Sigma, S, s_0, \delta, E, F)$ where:

- $\Sigma = \{\text{free}_\rho, \text{alloc}_\rho\}$, a subset of E
- $S = \{\text{allocated}_\rho, \text{freed}_\rho, \text{error}_\rho\}$
- $s_0 = \text{freed}_\rho$
- $\delta = \text{the relation } \{(\text{freed}_\rho, \text{alloc}_\rho, \text{allocated}_\rho), (\text{allocated}_\rho, \text{free}_\rho, \text{freed}_\rho), (\text{freed}_\rho, \text{free}_\rho, \text{error}_\rho), (\text{allocated}_\rho, \text{alloc}_\rho, \text{allocated}_\rho)\}$
- $E = \{\text{error}_\rho\}$
- $F = E$

An illustration of this monitor template can be seen in Figure 6.

Short-term Use-before-init Monitor Template




Figure 7: An illustration of a short-term use-before-init monitor template.

A use-before init monitor detects usage of a memory location before the location has been initialized, possibly leaving the resource in an unexpected state when it is accessed or used. Given a region ρ , a use-before-init monitor template is defined as the tuple $(\Sigma, S, s_0, \delta, E, F)$ where:

- $\Sigma = \{read_\rho, init_\rho\}$, a subset of E
- $S = \{uninitialized_\rho, initialized_\rho, error_\rho\}$
- $s_0 = uninitialized_\rho$
- $\delta = \text{the relation } \{(uninitialized_\rho, init_\rho, initialized_\rho), (uninitialized_\rho, read_\rho, error_\rho)\}$
- $E = \{error_\rho\}$
- $F = \{error_\rho, initialized_\rho\}$

An illustration of this monitor template can be seen in Figure 7.

3.2 Control Flow

EBA provides a representation of the control flow of the input source files which is utilized to detect bugs. EBA generates a tree structure of the input with each path in this tree structure modeling a possible execution path containing information about the modelled statements.

The control flow graph of a program can be seen as a finite state machine $(\Sigma, S, s_0, \delta, F)$, where Σ is the powerset of effects, S is a finite non-empty set of program points, s_0 is an element of S representing the entry point, δ is the state-transition function $\delta : S \times \Sigma \rightarrow S$ reflecting the edges of the control graph labeled by effects produced by computations and F is the empty set of final states.

The control flow generated by EBA is acyclic, since EBA unrolls loops within a fixed depth and generates a path of this length accordingly. I keep the abstract formulation since, in principle, the monitor automata checkers will work with more general abstractions over programs.

The powerset of effects, Σ , of the control flow abstraction is the set of all effects EBA detects, annotated by the region variables being affected by a given effect. S is the set of

program points after unrolling loops. The definition of the control flow abstraction is shown in the following, with a concrete example of a control flow formulated using this abstraction in Figure 8.

- $\Sigma = P\{\mathcal{E}_{\rho_i} | i \in \mathbb{N} \text{ and } \mathcal{E} \in \{alloc, free, read, write, uninit, call, lock, unlock\}\}$
- $S = \{1, 2, 3, 4, 5, 6, 7, 8\}$

The remainder of the automaton is defined according to the control flow being modelled, where the initial state, s_0 , is the entry point. A concrete definition of an example control flow is shown below with an accompanying illustration of this in Figure 8.

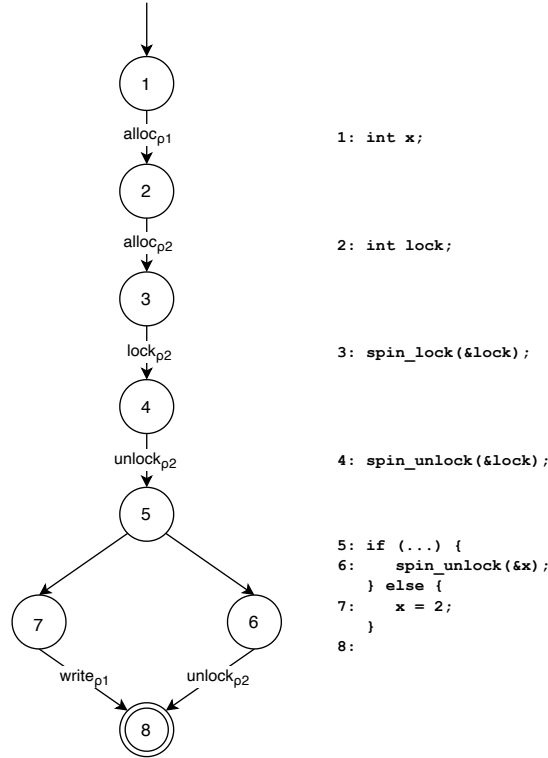


Figure 8: An illustration of a Control Flow automaton.

- $s_0 = 1$
- $\delta = \text{the relation } \{ \begin{array}{ll} (1, \{alloc_{\rho_1}\}, 2), & (2, \{alloc_{\rho_2}\}, 3), \\ (3, \{lock_{\rho_2}\}, 4), & (4, \{unlock_{\rho_2}\}, 5) \\ (5, \emptyset, 6), & (5, \emptyset, 7), \\ (6, \{unlock_{\rho_2}\}, 8), & (7, \{write_{\rho_1}\}, 8) \end{array} \}$
- $F = \emptyset$

Branches occur when an if-branch is encountered in the input source file and models the effects of the statements within if-statements.

To show the detection of a possible double-unlock bug we find the product of the control flow example shown in Figure 8 and the monitor generated from the template.

Given a region ρ , an automaton $A_{CFG} = (P(\mathcal{E}_{\bar{\rho}}), S, s_0, \delta, \emptyset)$ and a monitor template $A_{monitor} = (\sum, S', s'_0, \delta', F')$, the product automaton of A_{CFG} and $A_{monitor}$ is an automaton $P_{\bar{\rho}} = (\sum_{\bar{\rho}}, Q, s'_0, \delta_P, F_P)$ where

- $\mathcal{E}_{\bar{\rho}}$ is created from \mathcal{E} by substituting its template parameter with an actual value $\bar{\rho}$
- $F_{\rho} = S \times F'$
- $Q = S \times S'$
- $\delta_P : Q \times \sum \rightarrow Q$ which is generated for all $s_1, s_2 \in S, s'_1, s'_2 \in S', e \in \sum, E \in P(\mathcal{E}_{\bar{\rho}}) \forall q \in Q, q' \in Q'$ and $e \in \sum$ using the rules

$$\frac{s_1 \xrightarrow{\bar{e}} s_2 \quad s'_1 \xrightarrow{E} s'_2 \quad e \in E}{(s_1, s'_1) \rightarrow (s_2, s'_2)} \qquad \frac{s'_1 \xrightarrow{E} s'_2 \quad \forall e \in \sum . s_1 \not\xrightarrow{\bar{e}}}{(s_1, s'_1) \rightarrow (s_1, s'_2)}$$

where $\bar{e} = e[\bar{\rho}/\rho]$ for some $e \in \sum$ by substitution of the template parameter.

It is necessary to define rules for the state changes within this product, given that a monitor only accepts effects on a given region. E is the set of effects happening in a given program step, $s'_1 \rightarrow s'_2$. The state s_1 will change given that the transition happens on the effect e , present in E . If this is not the case, the control flow will have changed state to s'_2 , while the monitor has not and stays the same, s_1 .

In other words, the state of the automata should not change if the effect does not happen on the monitored region, but the automaton representing the control flow *should*. The observant reader might notice that if regions are no longer present since they go out of scope in a given program, only the rightmost of the two previous inference rules is relevant. This is an opportunity for an optimization in a bug checker, which is presently not exploited. Taking scopes into account would complicate the product construction significantly and the reader might appreciate not having to reason about more convoluted definitions. Lack of scope information does lead to the question of whether these "dangling" regions incur a significant performance cost when implemented. I will investigate and evaluate this in later sections.

The product of the control flow example shown in Figure 8 and a generated double-unlock monitor automaton can now be found in order to demonstrate that a possible bug is detected by the monitor, resulting in the following definition. This definition is illustrated in Figure 9.

- $\sum = \{Entry, Nil, (alloc, \rho), (free, \rho), (read, \rho), (write, \rho), (uninit, \rho), (call, \rho), (lock, \rho), (unlock, \rho)\}$
- $S = \{(allocated, \rho), (freed, \rho), (read, \rho), (written, \rho), (uninitialized, \rho), (called, \rho), ((locked, \rho), locked_{\rho}), ((unlocked, \rho), unlocked_{\rho}), ((unlocked, \rho), error_{\rho}), End\}$

- $s_0 = \text{Entry}$
- $\delta = \text{the relation } \{ (\text{Entry}, (\text{alloc}, \rho), (\text{allocated}, \rho)),$
 $(\text{allocated}, (\text{alloc}, \rho), (\text{allocated}, \rho)),$
 $((\text{allocated}, \rho), (\text{lock}, \rho), (\text{locked}, \rho)),$
 $((\text{locked}, \rho), (\text{unlock}, \rho), ((\text{unlocked}, \rho), \text{unlocked}_\rho)),$
 $((\text{unlocked}, \rho), \text{unlocked}_\rho), \text{unlock}, ((\text{unlocked}, \rho), \text{error}_\rho)),$
 $((\text{write}, \rho), \text{unlocked}_\rho), \text{write}, ((\text{written}, \rho), \text{unlocked}_\rho)),$
 $((\text{unlocked}, \rho), \text{error}_\rho), \text{Nil}, (\text{End}, \text{error}_\rho)),$
 $((\text{written}, \rho), \text{unlocked}_\rho), \text{Nil}, (\text{End}, \text{error}_\rho)) \}$
- $F = (\text{End}, \text{error}_\rho)$

This product is illustrated in Figure 9.

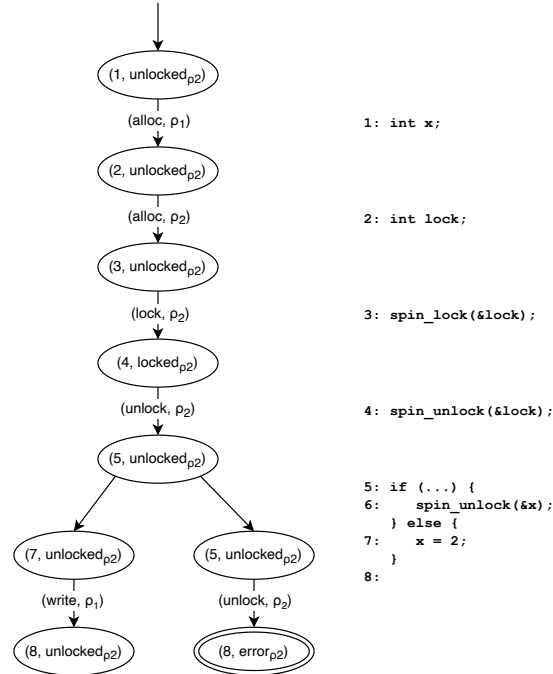


Figure 9: An illustration of the product construction of a double-unlock monitor automata and a control flow.

I have shown that it is possible to construct monitor templates which allows generating monitor automata monitoring effects happening on a region and shown that such a monitor can detect a possible bug in an example control flow. In order to implement this approach in practice, two things are needed, namely

- A control flow abstraction
- Concrete definitions of monitor templates

It is therefore necessary to implement these monitor templates, since the EBA framework can provide the control flow abstractions over a given input file. I will present the implementation of the monitor templates as part of this thesis in the following section.

4 Finding Double-Unlock Bugs in Practice

This section will detail how the abstractions defined previously are integrated into the EBA framework in order to explore the control flow graph of input programs while generating the product of the control flow and monitor automata in order to detect possible bugs.

4.1 Integration into EBA

In order to implement a new bug checker in EBA, the implementation has to conform to the structure set up by the framework of how these checkers should behave. Furthermore, signatures must be defined for these checkers in order to allow the framework to instantiate these checkers on the inferred regions of an input program. This section will describe how this has been accomplished.

The EBA framework allows specifying checker signatures whose implementations are executed on a given input source file. Checker signature implementations instantiate a given bug checker for a given bug type and the internal logic of the bug checker is run by the framework.

These bug checkers must conform to the existing signatures of EBA in order to allow the framework to instantiate a given bug checker after which an abstract representation of the input source file is passed to the instantiated checker.

In order to detect possible bugs in input programs, monitor automata are used to detect bugs over control flow graphs. The control flow graph is provided by the EBA framework, while the monitor automata have been defined and implemented as part of this thesis.

A signature for a checker which allows instantiation of monitor automata bug checkers has been defined as part of this thesis. This signature is then implemented in order to let EBA instantiate the implemented checker. A function, `check`, is the only requirement for implementing this signature and takes two parameters after which it returns a list of strings for each detected possible bug in the input source file as expected by the EBA framework. These parameters are the abstractions of the input file and each global function defined in this file, both of which are passed to the function by the framework. This mimics the implementation of the existing CTL checkers in EBA and allows for easy integration into the framework.

The aforementioned signature is implemented as a module, `Make`, which is used by EBA in order to run automata bug checkers conforming to an automata signature. Specifying a signature which all automata-based checkers must conform to ensures that the automata expose the required state and transition functions for them to run. This `Make` module expects an implementation of this `AutomataSpec` signature.

When inferred effects happening on regions are encountered in the control flow graph, a monitor is instantiated based on the monitor template defined in `Make`. The `check` function of the `Make` module explores the CFG provided by the EBA framework of the given file and applies the transition of the monitor using the effects of statements, represented as nodes in the CFG. The

control flow graph is represented as a tree-like structure, and this tree is then explored further until the end of each path in the tree is explored, resulting in a set of monitor states. This set of states can then be explored in order to determine if any monitors have reached accepting — Error — states. If such a state is present, a possible bug has been discovered. On-the-fly exploration is used since the design of EBA necessitates exploring effects when they are encountered due to performance reasons [1].

EBA can generate the required control flow abstraction, which is used for analysis. This thesis consists of integrating the implementation of monitor templates into the existing framework, using the control flow abstractions provided by EBA. In order to present how monitor templates are instantiated based on the given control flow, it is necessary to describe the control flow abstractions. EBA generates a tree structure of the input, modeling statements as so-called **steps**. A path in this tree structure models a possible execution path, with each **step** in a path containing information about the modelled statements. The implementation of **check**, initiating the generation of the control flow abstraction using EBA can be seen in Figure 10.

```

let check file declaration =
  let variable_info = Cil.(declaration.svar) in
  match variable_info.vstorage with
  | Static -> L.of_list []
  | _ ->
    let _, global_function = Option.get(AFile.find_fun
      file variable_info) in
    let path_tree = paths_of global_function in
    let results = explore_paths global_function path_tree
      Map.empty true in
    let states = Map.values results in
    let matches = Enum.fold (fun acc m ->
      (List.filter A.is_accepting m) @ acc) [] states in
    let matches_reversed = List.rev matches in
    let function_name = Cil.(variable_info.vname) in
    let pp = List.map (fun m -> A.pp_checker_state m function_name)
      matches_reversed in
    let pp_list = List.map (fun m -> PP.to_string m) pp in
    L.of_list pp_list

```

Figure 10: The implementation of **check** initializing the control flow generation using EBA.

Nodes in the CFG structure provided by EBA can be of one of four different types, each representing the input statement. Nodes representing if-statements in the source input result are *If*-nodes in the tree, containing two branches. If an If-node is discovered, the two branches from that node are explored and the union of the resulting states is found. Nodes representing the end of a branch are *Nil*-nodes in the tree. Nodes representing assumptions made after if-statements are either true or false are *Assume*-nodes, but are not used in this work since all branches are explored. Finally all other statements are represented as *Seq*-nodes, which contain information about the shapes and effects of statements. An illustration of these types can be seen in Figure 11. These *Seq*-nodes are of interest, since they allow analysis on effects.




Figure 11: An illustration of the CFG node types found in EBA.

Seq-nodes contain a *step* which models a statement in the input source code. When a *Seq*-node is discovered in the tree, the — possibly multiple — effects of its containing step are explored. An illustration of this node can be seen in Figure 13. These possibly multiple effects raise a problem; since a given step contains a set of effects, the order of these effects are therefore not known and all orders of executing these effects must be explored. This must be done since a given ordering of effects can lead to a bug, while a different order might not. The reader is encouraged to examine the example given in Figure 12 to understand why this is the case. In the example, two locks are taken on two different regions and these regions are then both unlocked in a called function. One of the regions is then unlocked again following the function call, leading to a double-unlock bug. If monitors are not instantiated for each region, an error state would occur when analyzing the effects of the function call even though the bug is in fact not found within the function, but just after the function.

```

int lock1;
int lock2;

void unlock_func()
{
    _spin_unlock(&lock1);
    _spin_unlock(&lock2);
}

int main()
{
    _spin_lock(&lock1);
    _spin_lock(&lock2);
    unlock_func();
    _spin_unlock(&lock2);
}

```

Figure 12: An example of multiple locks happening on different regions, leading to a bug on one region, but not the other.

All permutations of the set of effects must therefore be found and mapped to a given region, while also preserving the information of the other permutations for that given region. Furthermore, the transition function of the monitor automata must be evaluated on the current input, resulting in a new state of that automata which again must also be stored for that region.




Figure 13: An illustration of the step type and its containing structures found in EBA.

In order to accomplish this, the current state of the monitor monitoring a region needs to be copied and applied as the current state of the monitor for each permutation found for a given set of effects. This is done in order to find all possible states of the monitor when given all effect orderings as input. The resulting states are then stored in the map for the given region, and future effects on that region are then applied on these states. An illustration of this copying on two permutations followed by another effect happening on the same region can be seen in Figure 14, leading to a possible double-unlock bug.

Discuss correctness at meeting.




Figure 14: An illustration of how copies of a double-unlock monitor are instantiated given multiple permutations of effect orderings.

This map can be formalized as the function $m : \text{region} \rightarrow \text{checker_state}$ where *checker_state* is the internal state of the monitor automata. Using this map it is possible to apply each permutation of effects and fold this list of effects into a modified map with possibly altered automata states for their corresponding regions. The resulting map can be explored in order to find any accepting states of monitor automata for a given region.

Monitors are instantiated whenever a new region is encountered and are then stored in this map. If a given region is already present in the map, then the current state of the monitor(s) monitoring the region are applied on the transition function of the monitor(s). If the region is not already present in the map, a monitor is instantiated for each effect happening in the program point where this region was encountered. Given that the map maps a region to the state of monitor automatas, the length of the map will never be larger than the number of regions in the input source file. The size of the set of possible monitor automata states for a given region depends on the effects of a statement operating on a given region, though this can be improved by filtering effects, as we will see shortly.

Given a large number of possible effects of a statement the resulting set of permutations of these effects will naturally grow. A set of N effects will result in $N!$ permutations; in other words, the number of monitor states for a given region will therefore in the worst case be $|\text{effects}|!$. This N can be reduced by only applying effects to monitors which are of interest, namely effects which are in the alphabet, Σ , of the monitor template definition. Filtering effects caused by a program point to only include elements in Σ greatly reduces the number of monitors states, though this is inherently dependent on the template definition. Definitions with a restricted alphabet will perform better. Effects are only applied to the monitors monitoring a region if the alphabet of the monitor in question accepts this effect, in other words utilizing the product construction inference rules defined in Chapter 3.1.

The implementation of the exploration of control flow, finding states and adding these states to the map can be seen in Figure 15. This `explore_paths` function takes a control flow graph generated by EBA and an initially empty map as input parameters and results in a map of all inferred regions and the resulting monitor states for these regions. This is accomplished by recursively exploring the tree-like control flow graph and depending on the encountered program point, applies the possible effects of the program point on the transition function of the instantiated monitor, `A`. The resulting state(s) are then added to the map.

```

let rec explore_paths path map =
  let p = path() in
  match p with
  | Seq(step, remaining) ->
    let apply_transition effect map =
      let region = get_region effect in
      match region with
      | Some r ->
        let result = Map.find_default [A.initial_state] r map in
        let m = List.fold_left
          (fun acc s -> A.transition s effect step :: acc)
          [] result in
        Map.add r m map
      | None -> map
    in
    let input = EffectSet.filter is_in_transition_labels step.effs.may
      |> EffectSet.to_list in
    if List.is_empty input then explore_paths remaining map
    else
      (* Find all permutations of effects e.g. {{lock, unlock}
        -> {{lock, unlock}, {unlock, lock}} *)
      let permutations = permute input in
      let map = List.fold_left (fun map effects ->
        (* For each effect in a permutation, apply the transition,
          and add the result to the map. *)
        List.fold_left (fun map effect -> apply_transition effect map)
          map effects
        ) map permutations in
      explore_paths remaining map
  | Assume(_, _, remaining) ->
    explore_paths remaining map
  | If(true_path, false_path) ->
    let true_branch = explore_paths true_path map in
    let false_branch = explore_paths false_path map in
    Map.union true_branch false_branch
  | Nil -> map

```

Figure 15: The implementation of the EBA control flow while keeping track of states for regions. This function takes a control flow graph generated by EBA and an initially empty map as input parameters and results in a map of all inferred regions and the resulting monitor states for these regions.

When all paths in the CFG tree structure have been explored, the regions which map to accepting states along with their location and traces are extracted from the mapping and presented to the user as possible bugs.

This approach can be described in pseudocode as follows.

```

function EXPLORE_PATHS(tree_node, map)
  if tree_node is Nil then
    return map
  else if tree_node is If(t, f) then
    true_branch  $\leftarrow$  EXPLORE_PATHS(t)
    false_branch  $\leftarrow$  EXPLORE_PATHS(f)
    return UNION(true_branch, false_branch)
  else if tree_node is Seq(step, next) then
    effects  $\leftarrow$  step.effects
    if IS_EMPTY(effects) then return EXPLORE_PATHS(next, map)
    end if
    permutations  $\leftarrow$  FIND_PERMUTATIONS(effects)
    return
    FOLD(
      FOLD(APPLY_TRANSITION(effect, map), map, effects),
      map, effects)
       $\triangleright$  Find all possible states for permutations.
       $\triangleright$  Add these states to the region map.

    end if
  end function

function APPLY_TRANSITION(effect, map)
  region  $\leftarrow$  effect.region
  previous_states_for_region  $\leftarrow$  FIND_DEFAULT([initial_state], region, map)
  states  $\leftarrow$  FOLD(TRANSITION(state, effect, previous_states_for_region))
  return ADD(map, region, states)
end function

tree  $\leftarrow$  GENERATE_TREE(input_file)
map  $\leftarrow$  EXPLORE_PATHS(tree)

```

4.2 Monitor Signatures

The signature of monitor automata must be implemented in order to use the bug checker with EBA. The implementation of a given monitor automata is passed to the aforementioned **Make** module and is then used to evaluate states based on the effects of regions. The signature of the monitor automata specifies a **state** as a discriminated union type, describing the possible states of the automata as well as a transition function, **transition**, which takes a previous state of the monitor along with an input effect.

```

type state =
| Locked
| Unlocked
| Error of Effects.e

type checker_state = {
  current_state: state;
  trace: step list;
  matches: step list;
}

```

Figure 16: An illustration of the *checker_state* structure for a double-unlock monitor.

In order to provide the user with detailed error reports this state is encapsulated in a **checker state** structure which keeps track of the current trace through the CFG along with granular location details for discovered possible bugs. Providing this information requires that the current CFG node must also be passed to the automata, due to the architecture of the EBA framework. The full signature for the transition function is therefore $transition : checker_state \rightarrow effect \rightarrow step \rightarrow checker_state$. No references to instantiated monitor modules are stored in the map, merely the current states of monitors.

A concrete example of a *checker_state* structure can be seen in Figure 16. This transition function in other words operates on an effect which is part of the set of effect types and results in a new monitor state which is part of the set of possible states defined within the monitor, reflecting the definition of a monitor template seen in Section 3.1. A concrete example of an implementation of the transition function can be seen in Figure 17.

```

let transition previous input step =
  let next new_state = with_previous previous new_state step in
  let previous_state = previous.current_state in
  match previous_state with
  | Unlocked ->
    (match input with
    | Mem(Lock, _) -> next Locked
    | Mem(Unlock, _) -> next (Error input)
    | _ -> next previous_state
    )
  | Locked ->
    (match input with
    | Mem(Unlock, _) -> next Unlocked
    | _ -> next previous_state
    )
  | Error _ -> next previous_state

```

Figure 17: The implementation of the transition function of the double-unlock monitor template definition.

This implementation has been chosen for performance and simplicity reasons, since the transition

function of a monitor does not change after it has been defined and can therefore be implemented as a performant static function. Storing only previous states and taking this previous state into consideration in the implementation of transition functions results in a lower memory footprint and possibly higher readability of the implementation since an alternative implementation, e.g. based on dynamically assigning transitions to an internal map within the monitor, would increase the memory footprint and possibly make it hard to track what the possible states for a given monitor are for the reader.

As described previously, on-the-fly exploration has been used for ease-of-implementation and performance reasons. The exploration finds the product construction of the control flow graph and monitor on a step-by-step basis by application of the transition function of a monitor and the information stored in a program point in the control flow graph. The resulting pair of control flow and monitor state is found in the product construction of the control flow graph and monitor as described in Chapter 3.1. After finding this pair, the monitor state in the second element of this pair is, as mentioned previously, added to the region and state map in order to track monitor states for regions. The product construction is built step-by-step, but the complete product construction is not saved for later use, since the elements of interest in the resulting pair is the found monitor state. The inference rules defined in Chapter 3.1 for constructing the product, state that a transition — the exploration of the program points — will change the state of the control flow graph, but possibly leave the monitor state in its current state if the effects of a program point are not in the transitions of the monitor. This is reflected in the implementation described previously in this chapter, where monitors remain in their current state given an input that the monitor does not operate on. Therefore, even though the product construction is not stored or presented explicitly in the implementation, it is still constructed albeit step-by-step and is in a partial state on a given program point.

Consider the illustration in Chapter 3.1, with the relevant part of this illustration shown in Figure 18. When exploring the control flow graph, on encountering the effect `unlock` and region ρ , a monitor is instantiated for ρ in the unlocked state. If a given effect is not in the transitions of the monitor, as seen when encountering the effect `write`, the exploration continues onto the next program point, implicitly changing the state of the control flow graph by transitioning onto the next node in the graph, while leaving the monitor in its previous state as the inference rules in Chapter 3.1 state. When encountering an `unlock` effect, which is in the transitions of the monitor, the monitor will change state to an error state. This corresponds to the defined product construction, though instead of constructing the entire product, this partial product is found on-the-fly.




Figure 18: An illustration of the partial product construction of a double-unlock monitor automata and a control flow, based on Figure 9.

5 Evaluation

6 Future Work

7 Conclusion

References

- [1] Iago Abal, Claus Brabrand, and Andrzej Wasowski. Effective bug finding in c programs with shape and effect abstractions. In *VMCAI*, 2017.
- [2] Iago Abal, Claus Brabrand, and Andrzej Wasowski. Effective bug finding in c programs with shape and effect abstractions. pages 34–54, 01 2017.
- [3] IEEE and The Open Group. pthread_spin_unlock - unlock a spin lock object. <https://pubs.opengroup.org/onlinepubs/9699919799/>, 2017. Accessed: 2019-11-25.
- [4] George C Necula, Scott McPeak, Shree P Rahul, and Westley Weimer. Cil: Intermediate language and tools for analysis and transformation of c programs. In *International Conference on Compiler Construction*, pages 213–228. Springer, 2002.

8 Appendix